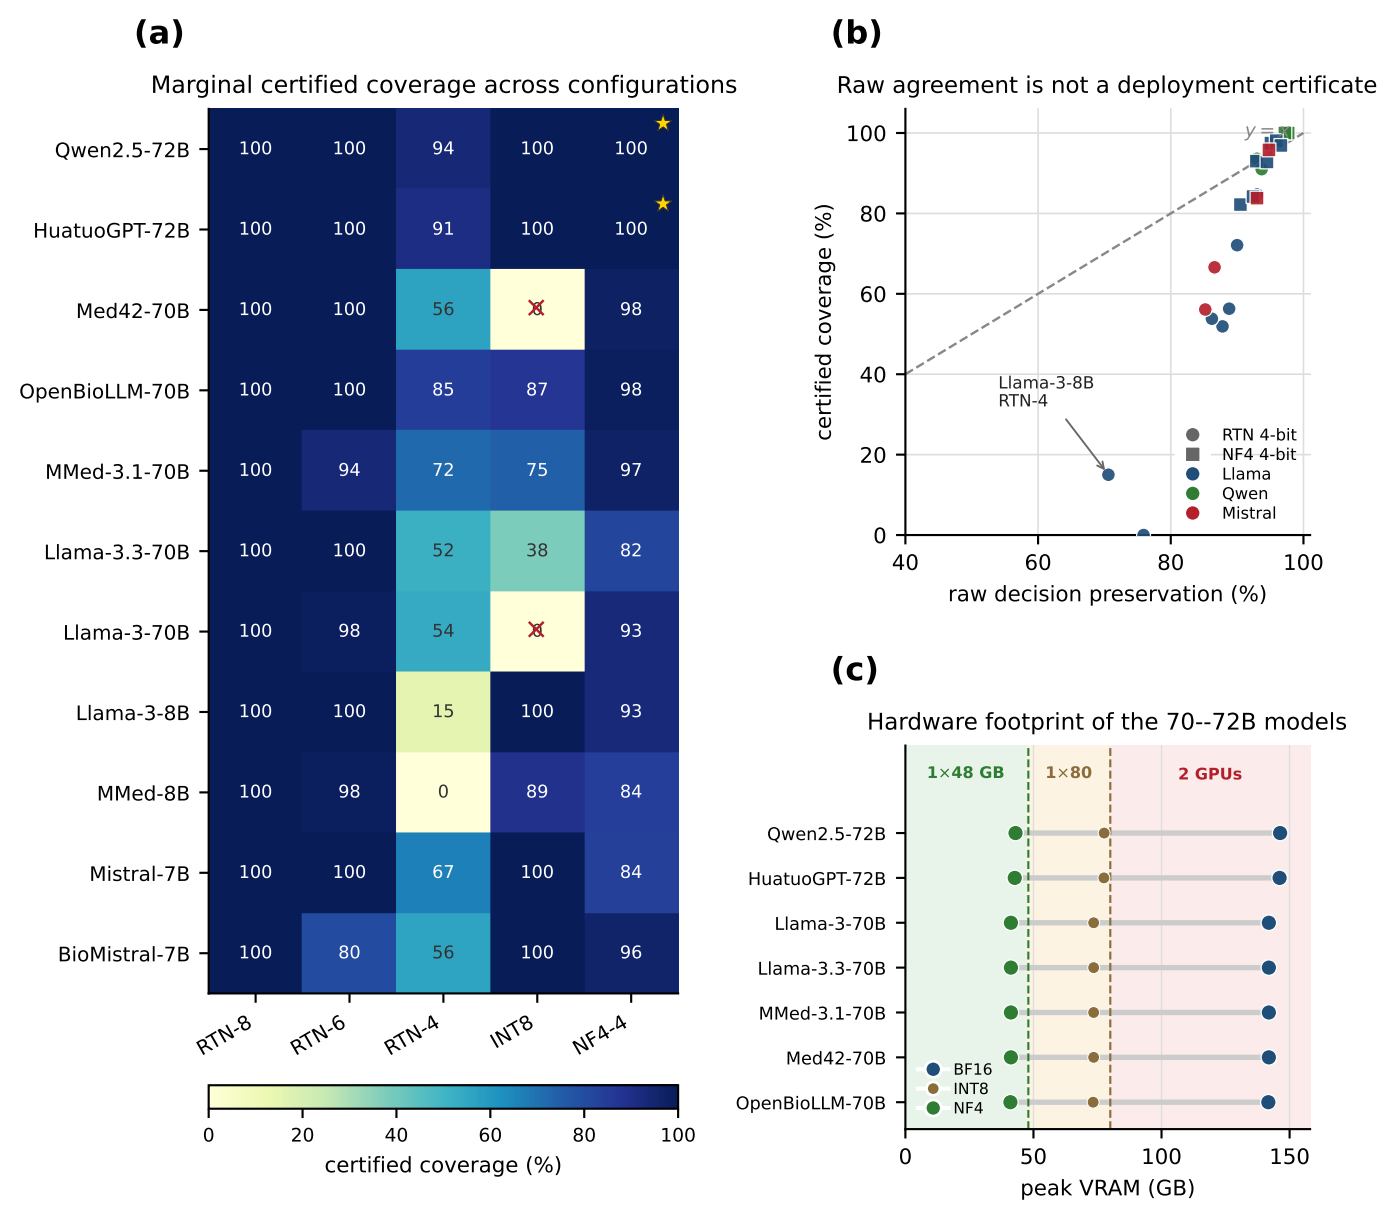
(a)
Marginal certified coverage across configurations
Qwen2.5-72B 100 100 94 100 100
HuatuoGPT-72B 100 100 91 100 100
Med42-70B 100 100 56 98
OpenBioLLM-70B 100 100 85 87 98
MMed-3.1-70B 100 94 72 75 97
Llama-3.3-70B 100 100 52 38 82
Llama-3-70B 100 98 54 93
Llama-3-8B 100 100 15 100 93
MMed-8B 100 98 0 89 84
Mistral-7B 100 100 67 100 84
BioMistral-7B 100 80 56 100 96
RTN-8
RTN-6
RTN-4
INT8
NF4-4
0 20 40 60 80 100
certified coverage (%)
(b)
Raw agreement is not a deployment certificate
certified coverage (%)
raw decision preservation (%)
Llama-3-8B
RTN-4
RTN 4-bit
NF4 4-bit
Llama
Qwen
Mistral
(c)
Hardware footprint of the 70--72B models
1×48 GB
1×80
2 GPUs
Qwen2.5-72B
HuatuoGPT-72B
Llama-3-70B
Llama-3.3-70B
MMed-3.1-70B
Med42-70B
OpenBioLLM-70B
BF16
INT8
NF4
peak VRAM (GB)