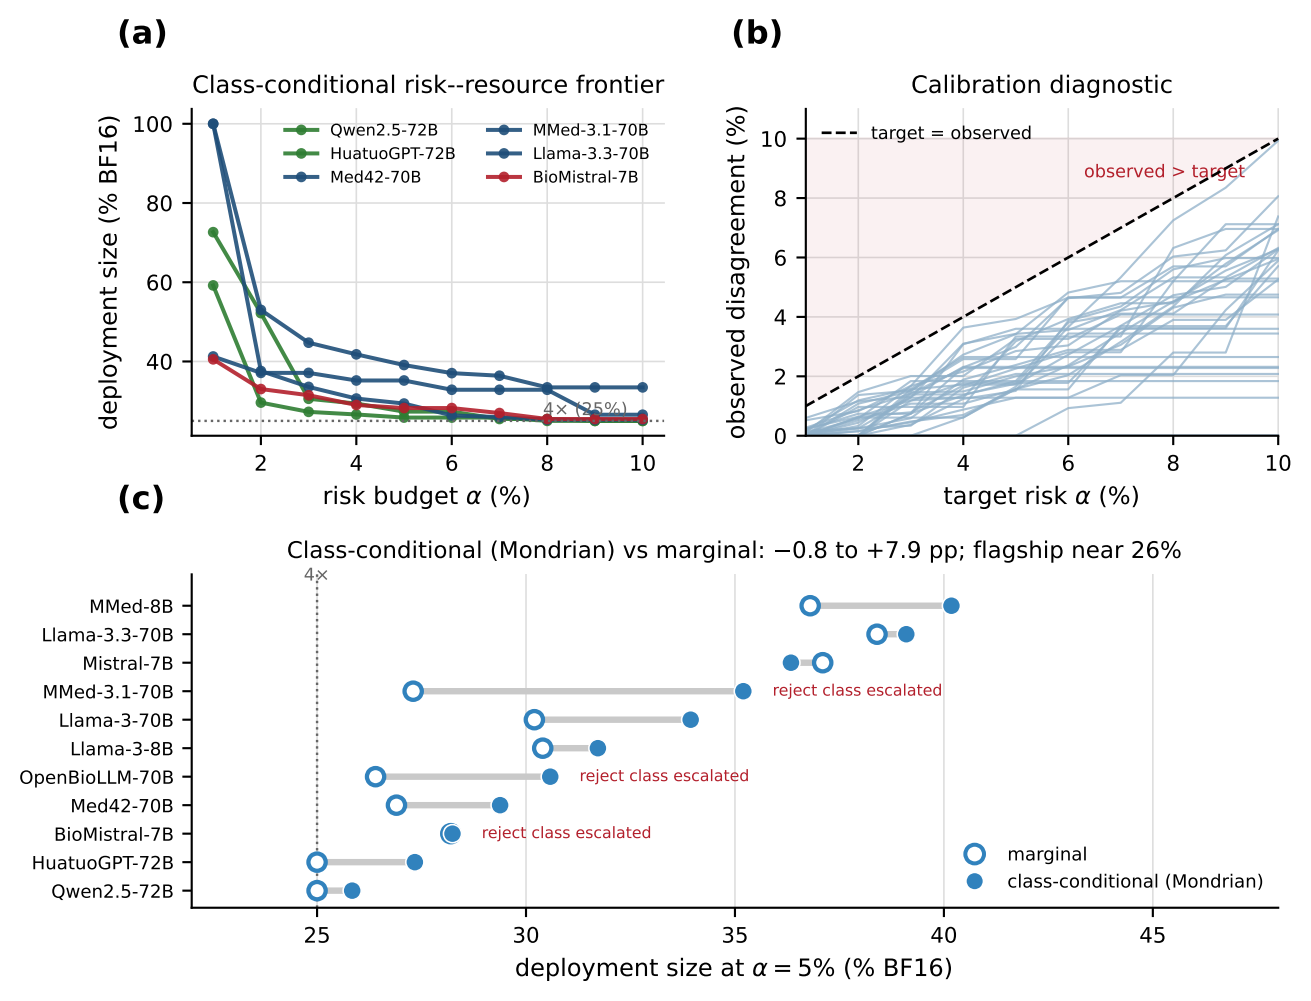
(a)
Class-conditional risk--resource frontier
deployment size (% BF16)
risk budget α (%)
Qwen2.5-72B
HuatuoGPT-72B
Med42-70B
MMed-3.1-70B
Llama-3.3-70B
BioMistral-7B
4× (25%)
(b)
Calibration diagnostic
observed disagreement (%)
target risk α (%)
target = observed
observed > target
(c)
Class-conditional (Mondrian) vs marginal: −0.8 to +7.9 pp; flagship near 26%
4×
MMed-8B
Llama-3.3-70B
Mistral-7B
MMed-3.1-70B
Llama-3-70B
Llama-3-8B
OpenBioLLM-70B
Med42-70B
BioMistral-7B
HuatuoGPT-72B
Qwen2.5-72B
reject class escalated
reject class escalated
reject class escalated
marginal
class-conditional (Mondrian)
deployment size at α = 5% (% BF16)